

From Technology Story to Public Issue

Modeling The Guardian's AI Coverage, 2021–2025

This project studies how AI-related coverage in *The Guardian* changed between 2021 and 2025. I use the Guardian Content API to construct a broad AI-related search corpus and a stricter AI-focused corpus. The final analysis focuses on the strict corpus because a manual audit shows that it has much higher precision than the broad-but-not-strict search results.

The main research question is:

How did *The Guardian's* AI-related coverage change between 2021 and 2025, and can article text and metadata predict whether an AI-related article is framed as a technology/business issue or as a broader public issue?

Main findings

- The expanded search strategy produced a large AI-related corpus across 60 months.
- The strict corpus is much more precise than the broad search corpus based on the manual audit.
- AI-related stories appear across many sections, not only Technology.
- Classification models use text, metadata, keyword indicators, and sentiment to predict whether an article belongs to a Technology/Business frame or a broader public-issue frame.
- LDA topic modeling identifies recurring themes in AI coverage, such as tools, business competition, regulation, education, copyright, labor, and risk.

Links

- [Interactive visualizations](#)
- [Full report](#)
- [Download the PDF report](#)
- [GitHub repository](#)